

SQL Injection Isn't Dead

Smuggling Queries at the Protocol Level

DEF CON 32 - August 10, 2024

Teaser

- In this talk, we will learn how...
 - ... queries travel from app to DB
 - ... attackers can inject there
 - ... prevalent the problem is


SELECT * FROM speakers

Name	Role	Company	Team
Paul Gerste	Vuln Researcher	Sonar	R&D

(1 row)

```
SELECT * FROM speakers, companies
```

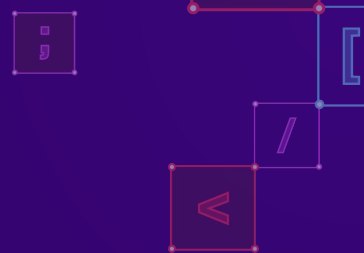
Name	Role	Company	Team
Paul Gerste	Vuln Researcher	Sonar	R&D

Logo	Name	Description
	Sonar	The home of Clean Code

(1 row)

Content

- The Idea
- Attacking Database Wire Protocols
- Real-World Applicability
- Future Research
- Takeaways

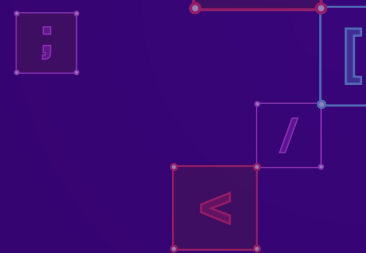


The Idea

Request smuggling, but for binary protocols

Prior Art

- James Kettle: HTTP Desync Attacks
 - Cause disagreement over the end of HTTP requests
- Most root causes:
 - Text parsing: 17 vs. \t17
 - Logical: CL vs. TE
- What about other protocols?



What About Binary Protocols?

- Binary protocols also need message boundaries
- How do they do it?
- Delimiters
 - E.g., null-terminated strings
- Length fields
 - E.g., Type-Length-Value (TLV) protocols

How To Desync Binary Protocols

- Delimiters
 - Insert delimiters into values

BLOG POST

Zimbra Email - Stealing Clear-Text Credentials via Memcache injection



Simon Scannell

VULNERABILITY RESEARCHER

June 14, 2022

DATE

How To Desync Binary Protocols

- Delimiters
 - Insert delimiters into values
- Length fields
 - 🤔
 - Endianness issues?
 - Overflows?

BLOG POST

Zimbra Email - Stealing Clear-Text Credentials via Memcache injection



Simon Scannell
VULNERABILITY RESEARCHER

June 14, 2022
DATE

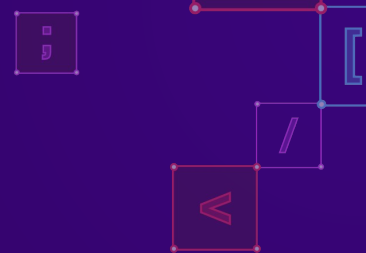
Where Are Binary Protocols Used?

- Databases
- Caches
- Message queues
- ... and many more

Attacking Database Wire Protocols

Why Database Wire Protocols?

- Extremely common
 - Almost every web app has a database
- Databases are high-value targets
 - Interesting data (e.g., PII)
 - Relevant data (e.g., for authentication)
- Guaranteed user input
 - Most queries contain some user input



High-Level Protocol Comparison

- PostgreSQL
- MySQL
- Redis
- MongoDB

High-Level Protocol Comparison

- **PostgreSQL**
- MySQL
- Redis
- MongoDB

Type	Length				Value...
'Q'	00	00	00	17	"SELECT ..."

High-Level Protocol Comparison

- PostgreSQL
- **MySQL**
- Redis
- MongoDB

Length			Sequence	Value...
00	00	17	00	"SELECT ..."

High-Level Protocol Comparison

- PostgreSQL
- MySQL
- **Redis**
- MongoDB

Type	Length	Delimiter	Value...	Delimiter
'+'	"17"	\r\n	"SELECT ..."	\r\n

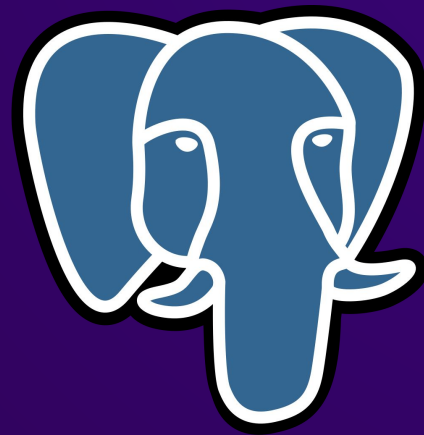
High-Level Protocol Comparison

- PostgreSQL
- MySQL
- Redis
- **MongoDB**

messageLength				requestID				responseTo			
17	00	00	00	00	00	00	00	00	00	00	00
opCode				value							
DD	07	00	00	...							

Case Study:

PostgreSQL



PostgreSQL Wire Protocol

Type	Length				Value...
'Q'	00	00	00	17	"SELECT ..."

- Type: 1-char identifier
- Length: 4-byte integer
- Value

PostgreSQL Wire Protocol

Type	Length				Value...
'Q'	00	00	00	17	"SELECT ..."

- Type: 1-char identifier
- Length: 4-byte integer
- Value

Max value: $2^{32}-1$

PostgreSQL Wire Protocol

Type	Length				Value...
'Q'	00	00	00	17	"SELECT ..."

- Type: 1-char identifier
- Length: 4-byte integer
- Value

Max value: $2^{32}-1$



The Bug

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst)  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

The Bug

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst)  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

● Write message type

The Bug

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst) —● Save size offset  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

The Bug

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst)  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

Write size

The Bug

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst)  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

The message buffer

The Bug

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst)  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

Buffer length (int)

The Bug

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst)  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

Truncate to int32

Message Size Overflow

Message 1					Message 2			
Type	Length				Value	Type	Length	
'Q'	00	00	00	08	"AAAA"	'Q'	00	00

Size: 8 = 0x00000008

4 bytes length + 4 bytes data

Payload: "A" * 4

Message Size Overflow

Message 1						Message 2			
Type	Length				Value	Type	Length		
'Q'	FF	FF	FF	FF	"AAAA..."	'Q'	00	00	

Size: $2^{32}-1 = 0x\text{FFFFFFFF}$

4 bytes length + $2^{32}-5$ bytes data

Payload: "A" * ($2^{32} - 5$)

Message Size Overflow

Message 1					?			
Type	Length				Value	?	?	
'Q'	00	00	00	04	""	'A'	'A'	'A'

Size: $2^{32} + 4 = 0x1000000004$

4 bytes length + 2^{32} bytes data

Payload: "A" * (2**32)

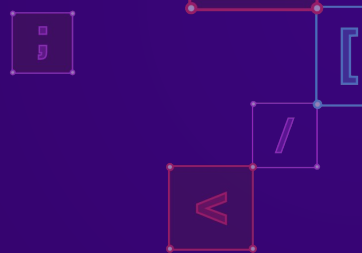
Message Size Overflow

Message 1						Injected Message			
Type	Length				Value	Type	Length		
'Q'	00	00	00	04	" "	'Q'	00	00	

Size: $2^{32} + 4 = 0x1000000004$

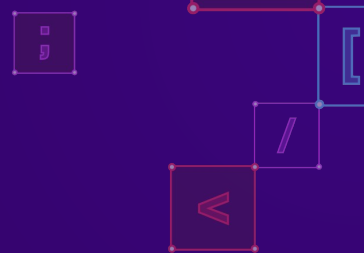
4 bytes length + 2^{32} bytes data

Payload: fakeMsg + "A" * (2**32 - len(fakeMsg))



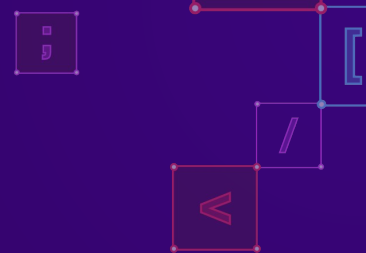
Crafting a Payload

- Simple payload is easy
 - `"INSERT INTO admins VALUES ... --".ljjust(2**32, "A")`
- But depends on the query
 - Where is the injection point?
 - How long is the query?
- Can we make it more reliable?
 - Yes!



Crafting a Payload

- Idea: NOP sled
 - Connection stays open on non-fatal errors
 - Spam a lot of small messages (5 bytes each)
 - Hit start of a message → success
 - Hit something else → connection closed
- Success after ≤ 5 attempts!
 - 20% chance of success
 - Attack is repeatable, just change the offset



Crafting a Payload

- Can we make it even better?
- NOP sled v2!
 - Overlapping pattern instead of tiny messages
 - Each byte should be a valid message start
- Constraints:
 - Max message size: `0x3fffffff` → first size byte cannot be `>0x3f`
 - No valid message type `≤0x3f` :/
- Solution: Each **2nd** byte should be a valid message start
 - Hit a valid type byte → **success**
 - Hit something else → **connection closed**
- Success after `≤2` attempts!
 - 50% chance of success
 - Attack is repeatable, just change the offset

Vulnerable Libraries

- Vulnerable:
 - Go: pgx, go-pg, pgdriver, [redacted]
 - C#: Npgsql
- Unexploitable:
 - Java: pgjdbc-ng, r2dbc-postgresql
 - JS: pg, pogi, postgres, @vercel/postgres

Exploitable Applications

- Any use one of the libraries is potentially exploitable
- Requirement: smuggle 4GB into a query
- Confirmed examples:
 - Mattermost (when file upload limit $\geq 4\text{GB}$)
 - Harbor!

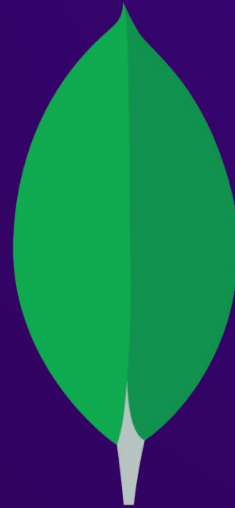
Demo: Harbor

- Container registry
- Used by VMWare Tanzu K8s
- Default configuration was vulnerable
- No authentication required



HARBOR

Case Study: MongoDB



MongoDB Wire Protocol

messageLength				requestID				responseTo			
17	00	00	00	00	00	00	00	00	00	00	00
opCode				value							
DD	07	00	00	...							

- Length field in the header
- Queries are BSON documents
 - Hierarchical objects
 - Serialized to TLV sections

The Bug

```
async fn write_to<T: AsyncWrite + Send + Unpin>(&self, mut writer: T) -> Result<()> {
    let sections = self.get_sections_bytes();
    let total_length = Header::LENGTH
        + std::mem::size_of::<u32>()
        + sections.len()
        + /* ... */;
    let header = Header {
        length: total_length as i32,
        // ...
    };
    header.write_to(&mut writer).await?;
    writer.write_u32_le(self.flags.bits()).await?;
    writer.write_all(&sections).await?;
    // ...
}
```

The Bug

```
async fn write_to<T: AsyncWrite + Send + Unpin>(&self, mut writer: T) -> Result<()> {
```

```
    let sections = self.get_sections_bytes();
```

```
    let total_length = Header::LENGTH
```

```
        + std::mem::size_of::<u32>()
```

```
        + sections.len()
```

```
        + /* ... */;
```

```
    let header = Header {
```

```
        length: total_length as i32,
```

```
        // ...
```

```
    };
```

```
    header.write_to(&mut writer).await?;
```

```
    writer.write_u32_le(self.flags.bits()).await?;
```

```
    writer.write_all(&sections).await?;
```

```
    // ...
```

```
}
```

• Get content bytes

The Bug

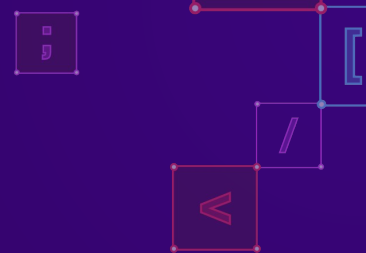
```
async fn write_to<T: AsyncWrite + Send + Unpin>(&self, mut writer: T) -> Result<()> {  
    let sections = self.get_sections_bytes();  
    let total_length = Header::LENGTH  
        + std::mem::size_of::<u32>()  
        + sections.len()  
        + /* ... */;  
    let header = Header {  
        length: total_length as i32,  
        // ...  
    };  
    header.write_to(&mut writer).await?;  
    writer.write_u32_le(self.flags.bits()).await?;  
    writer.write_all(&sections).await?;  
    // ...  
}
```

• Calculate message size (usize)

The Bug

```
async fn write_to<T: AsyncWrite + Send + Unpin>(&self, mut writer: T) -> Result<()> {
    let sections = self.get_sections_bytes();
    let total_length = Header::LENGTH
        + std::mem::size_of::<u32>()
        + sections.len()
        + /* ... */;
    let header = Header {
        length: total_length as i32,
        // ...
    };
    header.write_to(&mut writer).await?;
    writer.write_u32_le(self.flags.bits()).await?;
    writer.write_all(&sections).await?;
    // ...
}
```

• Truncate to i32



Crafting a Payload

- Must avoid bad bytes
 - Payload must be valid UTF-8
- Problem:
 - Message type (dd 07) is already invalid
 - Size fields can become invalid
- Solution:
 - Use metadata to create those bytes!

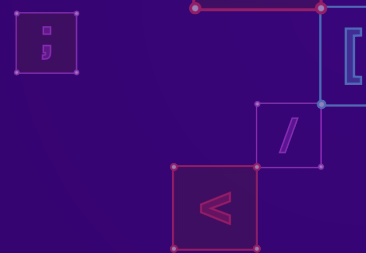
Demo: MongoDB Rust Driver

- The official MongoDB driver in Rust
- 3M downloads

Real-World Applicability

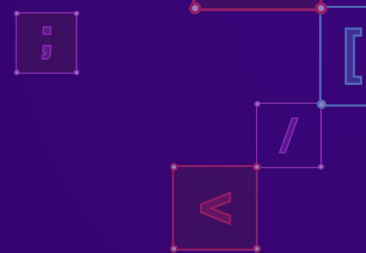
Constraints

- The elephant in the room:
 - Are 4GB of data realistic?
- Aren't apps limiting input sizes?
- Can \$language handle such big payloads?



How Web Apps Handle Large Payloads

- Aren't apps limiting input sizes?
- Common protections:
 - Default body size limits
 - Maximum JSON/form decode sizes
 - Size-limiting reverse proxies
 - ... and more



How Web Apps Handle Large Payloads

- Potential bypasses
 - Unprotected endpoints
 - Compression
 - WebSockets
 - Alternate body types
 - Incrementation

How Web Apps Handle Large Payloads

- Potential bypasses
 - **Unprotected endpoints**
 - Compression
 - WebSockets
 - Alternate body types
 - Incrementation
- Some have no default limits
- Some explicitly disable the limits
 - E.g., Harbor

How Web Apps Handle Large Payloads

- Potential bypasses
 - Unprotected endpoints
 - **Compression**
 - WebSockets
 - Alternate body types
 - Incrementation
- Some enforce size limits **before** decompression
 - E.g., Nginx

How Web Apps Handle Large Payloads

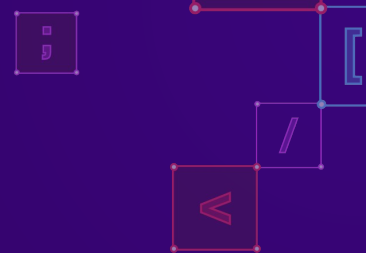
- Potential bypasses
 - Unprotected endpoints
 - Compression
 - **WebSockets**
 - Alternate body types
 - Incrementation
- Compression support
- Large message size
- Many filters don't apply

How Web Apps Handle Large Payloads

- Potential bypasses
 - Unprotected endpoints
 - Compression
 - WebSockets
 - **Alternate body types**
 - Incrementation
- Some filters don't apply
- E.g., multipart forms

How Web Apps Handle Large Payloads

- Potential bypasses
 - Unprotected endpoints
 - Compression
 - WebSockets
 - Alternate body types
 - **Incrementation**
- Concat/inflate strings on the server side
- Depends on the business logic



Language Comparison

- How well do languages handle big payloads?
 - How big can strings/buffers be?
 - Are integer overflows silent?

Language Comparison: Large Payloads

Language	Max. String Size	Max. Buffer Size
Go	$> 2^{32}$	$> 2^{32}$
Java	$2^{31}-1$	$2^{31}-1$
JS	$2^{29}-24 *$	$> 2^{32} *$
C#	$2^{31}-1$	$> 2^{32}$
Python	$> 2^{32}$	$> 2^{32}$
Rust	$> 2^{32}$	$> 2^{32}$

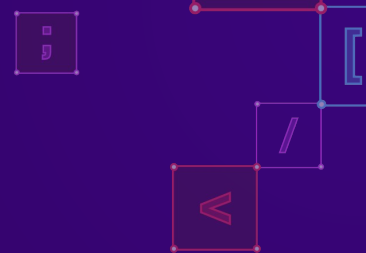
Only considering 64-bit versions.

* Depends on the implementation

Language Comparison: Integer Overflows

Language	Silent Addition Overflow?	Silent Serialization Overflow?
Go	Yes	No *
Java	Yes	No *
JS	No	Depends on impl.
C#	Yes	No *
Python	No	No
Rust	No	No *

* Type system prevents overflows



Real-World Applicability

- Can I send large payloads?
 - A lot of times, yes!
- Can integers silently overflow/truncate?
 - Sometimes!
- Can I exploit real-world apps with this?
 - Definitely!

Inspiration for Future Research

Safety First: No DoS Please!



Do not send large payloads to third-party systems!

Non-Invasive Detection

- White-box tests are harmless
 - Just set up your own test environment
- How to test this black-box?
 - Sending large payloads risks DoS
- More research and tools needed!
 - Can we safely detect vulnerable libraries?
 - Build tools to test this safely

Research More!

- More protocols
 - Other databases, caches, message queues
 - And more!
- Find more desync techniques
 - What about delimiters?
- More large payload methods
 - New ways to smuggle large payloads past defenses
 - Can we make the server generate the large payload?

Research More!

- All this was about **4-byte** length fields
- What about 2-byte fields?
 - Much easier to exploit (16KB vs. 4GB)
 - More to come in the future 🙄🙄

Conclusion

Takeaways

- Int overflows are relevant in memory-safe languages
- Sending large amounts of data is feasible
- If you can't hack it, just go a level deeper!

Questions?

@sonarsource
<https://sonarsource.com>

©2024, SonarSource S.A, Switzerland.

